

课程讲义

Learning to reason with LLMs

CS294/194-280: Advanced Large Language Model Agents

Jason Weston, Meta & NYU

中文教材化讲义 / Codex Harness Build



The graphic is a promotional poster for a MOOC. It has a dark blue background. On the left, there is a circular portrait of Jason Weston, a man with short blonde hair, wearing a colorful patterned shirt. Below the portrait, the text 'JASON WESTON' is written in white, and the Meta logo is below that. On the right, the text 'UC Berkeley Center for Responsible, Decentralized Intelligence' is in small white font. Below that, 'Advanced Large Language Model Agents MOOC' is written in large white font. Underneath, 'Learning to Reason with LLMs' is written in large white font. At the bottom right, 'February 3rd, 2025' is written in white. A decorative pattern of white dots is in the bottom right corner.

UC Berkeley Center for Responsible,
Decentralized Intelligence

Advanced Large Language
Model Agents MOOC

Learning to Reason
with LLMs

February 3rd, 2025

JASON
WESTON

Meta

课程页: <https://rdi.berkeley.edu/adv-llm-agents/sp25>

录播: https://www.youtube.com/live/_MNILhU33H0

slides: Jason-Weston-Reasoning-Alignment-Berkeley-Talk.pdf

补充 readings: DPO / IRPO / Chain-of-Verification

目录

1 本讲学习目标	2
2 背景与问题设置	2
2.1 把“学习推理”看成学习更好的 judge	2
2.2 System 1 失败为什么逼着我们走向 System 2	3
2.3 后训练谱系: SFT、RLHF、DPO	3
3 主体讲解	4
3.1 先把验证这件事做清楚: CoVe、System 2 Attention、Branch-Solve-Merge	4
3.1.1 Chain-of-Verification: 不是更长 CoT, 而是更可审计的 CoT	4
3.1.2 System 2 Attention 与 Branch-Solve-Merge	4
3.2 Self-Rewarding: 为什么 reasoning learning 会卡在 human bottleneck	5
3.2.1 标准 RLHF 的扩展瓶颈	5
3.2.2 Self-Rewarding LM: actor 和 judge 双能力模型	6
3.2.3 训练 recipe 与实验设计	7
3.3 IRPO: 把 preference optimization 从 style alignment 拉回 reasoning	8
3.3.1 为什么 reasoning 需要更强的 reward structure	8
3.3.2 IRPO 的边界条件	9
3.4 Thinking LLMs、Meta-Rewarding 与 EvalPlanner	9
3.4.1 Thinking LLMs: thought generation 不应只服务数学	9
3.4.2 Meta-Rewarding: judge 也要继续被训练	10
3.4.3 EvalPlanner: 让 judge 先计划再打分	11
4 关键公式、算法与 readings 的连接	11
4.1 DPO: 本讲所有 recipe 的优化执行器	11
4.2 IRPO: 从 response preference 走向 reasoning preference	11
4.3 CoVe: 验证先于自信	12
5 失败模式、边界条件与前后讲联系	12
5.1 本讲最重要的失败模式	12
5.2 与前后讲的联系	12
6 本章小结	13
7 复习题	13
8 深入思考题	14
9 延伸阅读	14

1 本讲学习目标

第二讲把课程焦点从“推理时怎么花预算”转向“如何把 reasoning 行为直接教给模型”。Jason Weston 的核心论点是：如果我们只把语言模型当成会模仿文本的 System 1，那么它会持续暴露出 hallucination、sycophancy、jailbreaking 等模式匹配型失败；真正困难的地方在于，**如何让模型既会回答，也会评价回答，从而形成自我改进的闭环。**

读完本讲后，读者应能回答：

- 为什么 Weston 把 reasoning learning 的问题重写成“如何得到更可靠的 judgments / rewards”。
- SFT、RLHF、DPO 在 post-training 谱系中各自扮演什么角色，为什么 DPO 会成为讲者首选的训练积木。
- CoVe、System 2 Attention、Branch-Solve-Merge 这些方法如何把“长一点的思考”改写成“显式验证与分解”。
- Self-Rewarding LM、IRPO、Thinking LLMs、Meta-Rewarding、EvalPlanner 之间的递进关系是什么。
- 为什么 reasoning 训练最难的往往不是 actor，而是 evaluator / judge。

2 背景与问题设置

2.1 把“学习推理”看成学习更好的 judge

Weston 一开场就把目标说得很尖锐：理想 AI 应该尽可能自己给自己出题、自己判断做得好不好、再根据这些判断继续更新自己。这不是简单自动化，而是在追问一个更基础的问题：当模型能力越来越接近甚至超过普通人时，谁来继续稳定地给它提供训练信号？

Goal: An AI that "trains" itself as much as possible

- Creates new tasks to train on (challenges itself)
- Evaluates whether it gets them right ("self-rewarding")
- Updates itself based on what it understood

Research question: *can this help it become superhuman?*




图 1: Weston 对 self-training AI 的定义：能生成任务、分配 reward、并把这些 judgments 重新喂回训练流程。

这一定义已经隐含了本讲的中心思想。Reasoning 不仅是“生成更长的 chain-of-thought”，而是：模型能否获得足够高质量的反馈信号，判断哪条思路更好、哪条思路只是看起来像思考。

2.2 System 1 失败为什么逼着我们走向 System 2

在 Weston 的叙述里，当前 LLM 更像 **System 1**：反应快、擅长联想、每个 token 都是固定计算，但很容易学到不该学的相关性。幻觉、迎合用户错误观点、越狱，本质上都是“反应式模式匹配”超出安全工作区后的副产物。换句话说，问题不只是模型会错，而是它经常没有显式的机制停下来检查自己。

Weston 因此把 reasoning 引向了一个很务实的方向：与其要求模型凭空更聪明，不如先让它学会更系统地**验证 (verification)**、**重写 (rewriting)**、**分解 (decomposition)** 和**评价 (judging)**。这就是他所说的 System 2 路线。

2.3 后训练谱系：SFT、RLHF、DPO

在进入 reasoning-specific 方法之前，讲者先回顾 pre-o1/r1 时代的 post-training 谱系：SFT 用任务示例继续做 next-token learning，RLHF 引入人类偏好与 reward model，而 DPO 则把 RLHF 的目标改写成一个更稳定直接的 preference optimization loss。




图 2: SFT、RLHF、DPO 构成 Weston 后续所有训练 recipe 的基本骨架。

DPO 的标准写法是：

$$\mathcal{L}_{\text{DPO}}(\theta) = -\mathbb{E}_{(x, y_w, y_l)} \left[\log \sigma \left(\beta \log \frac{\pi_{\theta}(y_w | x)}{\pi_{\text{ref}}(y_w | x)} - \beta \log \frac{\pi_{\theta}(y_l | x)}{\pi_{\text{ref}}(y_l | x)} \right) \right]$$

其中 π_{θ} 是要训练的模型， π_{ref} 是参考模型， x 是输入指令， y_w 与 y_l 分别是胜出和败北回答， β 控制偏离参考模型的强度。直觉上，DPO 不再单独训练一个 reward model 再跑 RL，而是直接把“更喜欢 y_w 而不是 y_l ”写进模型更新目标。

这条公式为什么重要

Weston 之后讲的 self-rewarding、IRPO、meta-rewarding 都要回到同一个问题：偏好对从哪来？DPO 本身不是答案，它只是一个非常好用的“执行器”。真正困难的是如何构造足够可信的 (y_w, y_l) 。

3 主体讲解

3.1 先把验证这件事做清楚：CoVe、System 2 Attention、Branch-Solve-Merge

3.1.1 Chain-of-Verification：不是更长 CoT，而是更可审计的 CoT

CoVe 的贡献在于，它把“再想一遍”拆成四个可解释的步骤：先给草稿答案，再生成 verification questions，然后独立回答这些问题，最后整合成最终答案。

Chain-of Verification Reduces Hallucination in Large Language Models

• Shehzaad Dhuliawala, Mojtaba Komeili, Jing Xu, Roberta Raileanu, Xian Li, Asli Celikyilmaz, Jason Weston




图 3: CoVe 把验证流程显式拆解，核心目的是减少最终回答对原始错误草稿的路径依赖。

```
1 Draft an initial answer
2 Plan verification questions
3 Answer each verification question independently
4 Synthesize a final verified response
```

这和第一讲里那种“让模型多采样几条 CoT 再投票”不同。CoVe 面向的是 factuality / hallucination 场景：如果错误来自事实断言，而不是搜索宽度不足，那么更关键的是把事实核查环节显式写进流程。

为什么要独立回答 verification questions 如果 verification step 继续强依赖原始草稿，那么模型只是在原错误附近做润色。CoVe 要求各验证问题独立作答，目的正是防止 answer contamination。

3.1.2 System 2 Attention 与 Branch-Solve-Merge

Weston 随后给出另外两类 System 2 方法。System 2 Attention (S2A) 不急着重写问题，尽量移除与任务无关的偏见和干扰项，再基于重写后的问题作答。

System 2 Attention (S2A)

Jason Weston, Sainbayar Sukhbaatar

Decide what to attend explicitly (system 2) by rewriting the input

Problem: whole context affects LLM output *even irrelevant parts!*

Hypothesis: soft-attention inherently spreads attention thin over everything. Also, LM objective favors correlations.

Solution: Make attention more explicit & effortful → **Prompt LLM to extract relevant context**

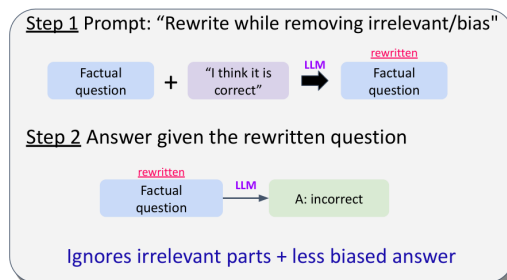


图 4: S2A 的关键是先决定“应该注意什么”，而不是直接在有偏输入上给答案。

这背后的直觉非常实用：不少 hallucination 不是知识完全缺失，而是模型注意力被噪声、用户偏见或表述方式带偏了。S2A 先做 **attention selection**，再做 answering。

Branch-Solve-Merge 则把复杂评价或生成任务拆成多个子问题，分别求解后再融合。它更像“评估任务也需要 planning”。和第一讲的 Tree of Thoughts 很像，只是这里的对象更多是 **response evaluation** 而不是 search over candidate answers。

本节的统一主线

Weston 给出的 System 2 方法有一个共同点：它们都把“多花一点 token”变成了“对哪一步花 token、为什么花 token”的结构化决策。

3.2 Self-Rewarding: 为什么 reasoning learning 会卡在 human bottleneck

3.2.1 标准 RLHF 的扩展瓶颈

当模型还不强时，人类给偏好标签并不太困难；一旦回答很长、主题专业、甚至模型在某些方面超过普通评审，人类就会越来越难稳定判断“哪个更好”。

Standard RLHF alignment approach: *use humans in the loop*

- first to create (X, Y) data;
- then to collect judgments on (X, Y') data

Image from: <https://arxiv.org/pdf/2009.01325.pdf>

图 5: 标准 RLHF 把人类放在数据生产和偏好判断两个关键位置, 因此扩展能力直接受限于人类评审吞吐量。

Weston 的研究问题很明确: 如果继续提升模型需要更高质量的 judgments, 而 humans in the loop 已经成为瓶颈, 那么能不能让模型学习生成这些 judgments ?

3.2.2 Self-Rewarding LM: actor 和 judge 双能力模型

Self-rewarding LM 的定义不是“模型给自己打分”这么简单。Weston 强调它至少要有两种能力:

- **instruction following capability**: 能对用户指令给出较好回答。
- **evaluation capability**: 能区分更好和更差的回答, 并给出可用的偏好信号。

Our approach

- Self-rewarding LMs
 - **Key idea**: train a self-rewarding language model that
 - 1) Has instruction following capability, i.e., given a user instruction, can respond to it appropriately

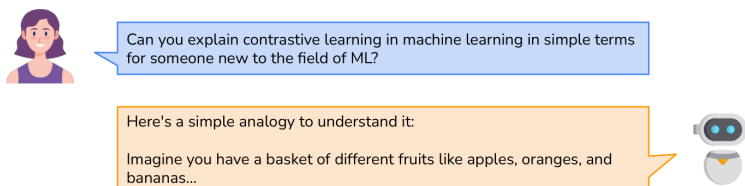


图 6: Self-rewarding 模型既是 actor 也是 judge。Weston 的关键不是“省人力”, 而是让 evaluator 能随着 actor 一起成长。

可以用一个极简抽象写成:

$$\mathcal{D}_{t+1} = \mathcal{D}_t \cup \{(x, \{(y_i, r_i)\}_{i=1}^k)\}$$

其中 $\mathcal{D}_t$ 是已有数据集，模型对同一指令 x 生成多个回答 y_i 并赋予分数 r_i ，然后把这些 judged candidates 再变成新的训练材料。

为什么这不是自嗨 如果模型只会生成、不会评估，那么新的训练数据不过是旧偏差的放大器。Self-rewarding 的真正难点在于 judge 能否提供**比现有 actor 更可靠的排序**。

3.2.3 训练 recipe 与实验设计

Lecture 给出的 recipe 很具体：从一个带少量 seed instruction-following data 和 evaluation data 的模型出发，在每轮中生成 prompts、candidate responses 和 self-rewards，然后把选出的 preference pairs 用 DPO 继续训练。


Our approach

Recipe 🍳: LM finetuned on small seed data.

Iterate 2 steps:

- (1) Self-instruction creation: generate prompts, responses & self-rewards with LM
- (2) Instruction-training: Train (DPO) on selected preference pairs

Iterations improve instruction following & reward modeling ability!



Self-Rewarding Language Models. Our self-alignment method consists of two steps: (i) *Self-Instruction creation*: newly created prompts are used to generate candidate responses

图 7: Self-rewarding 训练 recipe: 生成、判断、筛选、DPO，再迭代。

```

1 Initialize M1 with seed instruction-following and evaluation data
2 For each iteration t:
3     generate prompts, candidate responses, and self-rewards with M_t
4     select preference pairs from the judged candidates
5     run DPO to obtain M(t+1)

```

讲者特别展示了 LLM-as-a-Judge prompt 的维度设计：relevance、coverage、usefulness、clarity、expertise。这个细节非常关键，因为 reasoning training 不是只要一个 scalar reward 就够了；如果 reward 的来源不可解释、不可审计，那么后续迭代只会更脆弱。

Experiments

- LLM-as-a-Judge prompt
 - Instructs the LLM to evaluate the response using five additive criteria (relevance, coverage, usefulness, clarity and expertise)
 - Performs better than multiple choice format prompt

Review the user's question and the corresponding response using the additive 5-point scoring system described below. Points are accumulated based on the satisfaction of each criterion:

- Add 1 point if the response is relevant and provides some information related to the user's inquiry, even if it is incomplete or contains some irrelevant content.
- Add another point if the response addresses a substantial portion of the user's question, but does not completely resolve the query or provide a direct answer.
- Award a third point if the response answers the basic elements of the user's question in a useful way, regardless of whether it seems to have been written by an AI Assistant or if it has elements typically found in blogs or search results.
- Grant a fourth point if the response is clearly written from an AI Assistant's perspective, addressing the user's question directly and comprehensively, and is well-organized and helpful, even if there is slight room for improvement in clarity, conciseness or focus.
- Bestow a fifth point for a response that is impeccably tailored to the user's question by an AI Assistant, without extraneous information, reflecting expert knowledge, and demonstrating a high-quality, engaging, and insightful answer.

User: <INSTRUCTION_HERE>

<response><RESPONSE_HERE></response>

After examining the user's instruction and the response:

- Briefly justify your total score, up to 100 words.
- Conclude with the score using the format: "Score: <total points>"

Remember to assess from the AI Assistant perspective, utilizing web search knowledge as necessary. To evaluate the response in alignment with this additive scoring model, we'll systematically attribute points based on the outlined criteria.

图 8: Weston 强调 judge prompt 要拆解成多个 criteria, 避免 evaluator 只学到模糊的“我喜欢这个回答”。

实验结果的真正含义 Self-rewarding 模型在 instruction following 和 reward modeling 两个轴上都持续提升, 这说明“学会 judging”不会天然拖累“学会 acting”。但讲者没有把这个结果夸成万能方案, 因为 reasoning tasks 仍然比一般 instruction following 更难, 需要更强的可验证信号。

3.3 IRPO: 把 preference optimization 从 style alignment 拉回 reasoning

3.3.1 为什么 reasoning 需要更强的 reward structure

Weston 在第一个 self-rewarding 结果之后立刻承认一个限制: 模型确实更会回答、也更会评分了, 但 reasoning task 仍然提升有限。原因是通用偏好标签更像 style / helpfulness alignment, 而 reasoning task 需要对中间思路的成败有更细粒度的 credit assignment。

Iterative reasoning preference optimization

Richard, Weizhe, Cho, He, Sainaa, Jason

2024 (April)

Goal: use same self-rewarding type techniques, but for reasoning tasks..

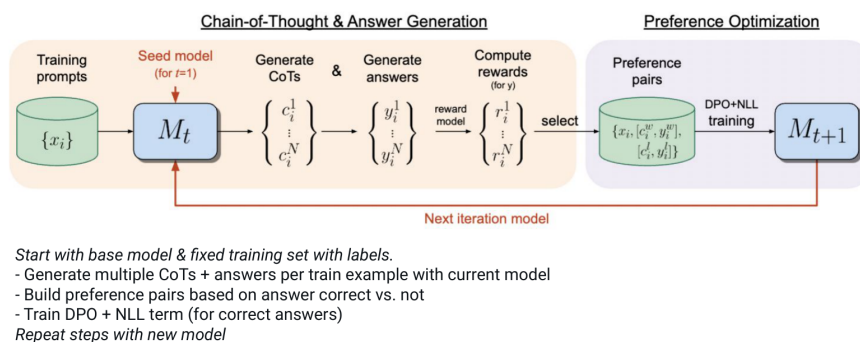


图 9: IRPO 的出发点: reasoning task 需要的不只是好回答, 而是“哪条思路最终导致了可验证正确答案”。

IRPO 的基本流程是: 对每个问题采样多条 chain-of-thought, 抽取最终答案, 用可验证规则判断

最终答案是否正确，再保留那些通向正确 final answer 的 reasoning traces，把它们与失败 traces 组成 preference pairs。

它的目标写成：

$$\mathcal{L}_{\text{IRPO}} = \mathcal{L}_{\text{DPO}} + \lambda \mathcal{L}_{\text{NLL}}(y_{\text{cot}}^*)$$

其中 $\mathcal{L}_{\text{DPO}}$ 负责学习 winning / losing CoT 的相对偏好， $\mathcal{L}_{\text{NLL}}$ 则继续鼓励模型模仿正确 reasoning trace， y_{cot}^* 表示到达可验证正确 final answer 的那条 CoT。

```
1 Sample multiple CoT candidates per problem
2 Extract the final answer from each candidate
3 Keep trajectories whose final answer is verifiably correct
4 Create winning/losing reasoning pairs
5 Train with DPO plus an NLL term on the winning traces
```

为什么 negative examples 是关键 Lecture 明确指出：如果只做 SFT, chosen 和 rejected generations 在 token 级概率上常常被拉得太近，模型无法真正学会“哪种 reasoning style 会通往正确答案”。IRPO 需要明确失败样本，才能把 reasoning 区分开来。

3.3.2 IRPO 的边界条件

IRPO 最大的优点是 reward 来自**可验证 final answer**，这让 reasoning learning 有了更坚实的地基。但它也有明确边界：

- 任务必须存在可验证答案，至少 final answer 要能自动判对错。
- 对开放式任务、创造性任务、模糊评价任务，IRPO 的 reward grounding 会变弱。
- 它仍然是在固定问题集上循环，不等于 agent 已经能在开放环境中自发发现新 reasoning tasks。

3.4 Thinking LLMs、Meta-Rewarding 与 EvalPlanner

3.4.1 Thinking LLMs: thought generation 不应只服务数学

Weston 进一步指出，o1 / R1 式“先想后答”不应该只限于数学或可验证推理。Thinking LLMs / Thought Preference Optimization (TPO) 的目标，是让模型在一般 instruction following 任务里也学会 thought generation。

Thinking LLMs: General Instruction Following with Thought Generation

Tianhao Wu, Janice Lan, Weizhe Yuan, Jiantao Jiao, Jason Weston, Sainbayar Sukhbaatar

2024 (October)

Trains LLMs to think & respond for "all" instruction following tasks, not just math

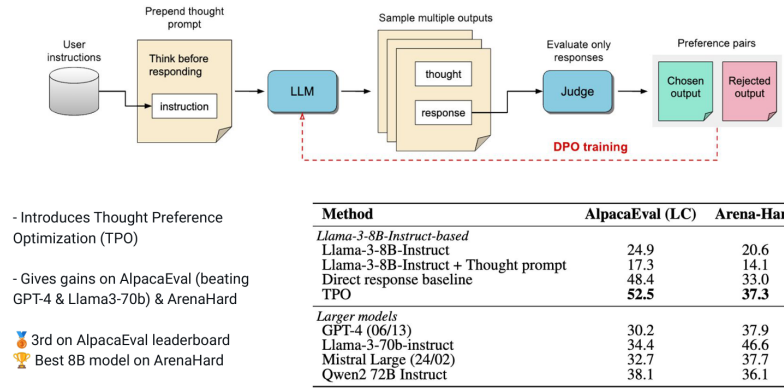


图 10: Thinking LLMs 的主张: thought generation 应该成为更广泛 instruction following 的默认能力, 而不是 reasoning benchmark 的特例。

这里的关键不是“每道题都写一长串想法”, 而是让模型学会何时需要显式计划、何时只需快速回答。对 agent 来说, 这意味着 reasoning policy 也应该被训练, 而不仅是 final answer style。

3.4.2 Meta-Rewarding: judge 也要继续被训练

如果 self-rewarding 把 actor 和 judge 放进同一个模型, 下一步自然是: judge 本身怎么变强? Meta-Rewarding 的回答是, 让模型不仅学习回答, 还学习对 judgment 进行 meta-judgment。

Recipe 🍳:

Iterate 3 steps:

- (1) Create Actor data: generate responses & self-rewards (judgments) with LM
- (2) Create Judge data: generate meta-rewards over judgments with LLM-as-a-Meta-Judge
- (3) Train DPO on preference pairs to both learn to act (1) AND to judge (2)




图 11: Meta-Rewarding 的三步循环: 生成 actor data、生成 judge data、同时训练 actor 与 judge。

一个简化的写法是:

$$p(i \succ j) = \sigma(s_i - s_j)$$

其中 s_i 和 s_j 是两个 judgments 的 latent score, σ 把分差转成偏好概率。Lecture 用 Elo-style intuition 说明: 如果能比较 judgment 本身谁更好, 我们就能反过来训练更强的 judge。

- 2 Create judge data: meta-judge comparisons over those judgments
- 3 Train DPO objectives for both the actor and the judge

Meta-Rewarding 的价值在于，它直面了“judge 会不会比 actor 落后”的问题。许多系统失败，不是因为生成器不会说，而是因为筛选器不会选。Weston 的路线是把筛选器也当成要持续学习的对象。

3.4.3 EvalPlanner: 让 judge 先计划再打分

最后，EvalPlanner 把同样的想法推到 evaluation task 上：若评估本身也很难，那么 judge 也应该有自己的 chain-of-thought、planning 和可验证训练数据。

Learning to Plan & Reason for Evaluation with Thinking-LLM-as-a-Judge

Swarnadeep Saha, Xian Li, Marjan Ghazvininejad, Jason Weston, Tianlu Wang
FAIR at Meta

We can also push reasoning further for the evaluation task.

EvalPlanner – a method to train o1/r1-like chain-of-thought (CoT) for the evaluation / reward model task.

This "Thinking-LLM-as-a-Judge" learns to generate planning & reasoning CoTs for evaluation.

Example Input and Output of EvalPlanner

User Instruction: [Text] There is a French pub called The Phoenix that has a low customer rating but is child friendly and is more than 30 dollars and is located in the riverside area near Crowne Plaza Hotel.
[Question] According to the above [Text], extract the attributes (name, eatType) of the restaurant.
Do not output your analysis or the result in the format of "attribute1[*value1*], attribute2[*value2*] ..."

Response A: Name[The Phoenix], eatType[pub], location[Riverside area near Crowne Plaza Hotel], priceRange[More than 30 dollars]

Response B: Name[The Phoenix], eatType[pub]

Output:
[Start of Evaluation Plan]
To evaluate the response quality, we will follow these steps:
1. "Check if the response provides the correct attributes". Verify that the response includes the attributes "name" and "eatType" as requested in the question.
2. "Verify the accuracy of the attribute values". Check if the values provided for the attributes match the information given in the original text.
3. "Check for any additional or missing information". Ensure that the response does not include any extra attributes or information not requested in the question.
[End of Evaluation Plan]
Step-by-step reference answer:
1. The response should include the attributes "name" and "eatType".
2. The values for the attributes should be:
* name: The Phoenix
* eatType: French pub
3. The response should not include any additional attributes or information.
[End of Evaluation Plan]
[Start of Plan Execution]
Step 1: Check if the response provides the correct attributes
Assistant A's response includes the attributes "Name", "eatType", "location", and "priceRange". However, the question only requests the attributes "name" and "eatType". Assistant B's response includes the attributes "name" and "eatType", which matches the request.
Step 2: Verify the accuracy of the attribute values
Assistant A's response provides the value "The Phoenix" for the attribute "Name" and "pub" for the attribute "eatType". However, the original text describes the pub as "French", which is missing in the response. Assistant B's response provides the value "The Phoenix" for the attribute "name" and "pub" for the attribute "eatType". Similar to Assistant A's response, it is missing the "French" descriptor.
Step 3: Check for any additional or missing information
Assistant A's response includes additional attributes "location" and "priceRange" that are not requested in the question. Assistant B's response does not include any additional attributes or information.
Based on the evaluation, both responses have some inaccuracies. However, Assistant B's response is more accurate in terms of providing only the requested attributes, even though it is missing the "French" descriptor. The final verdict is: [B]
[End of Plan Execution]
[Final Verdict]
[B]

图 12: EvalPlanner 把 evaluation task 转成可验证推理任务，训练的是“会规划的 judge”。

- 1 Generate a good response y to prompt x
- 2 Perturb x into a similar prompt x' and generate y'
- 3 Convert the pair into a verifiable evaluation task
- 4 Train a thinking judge to plan before scoring

这一步非常关键。到这里，本讲已经把 reasoning learning 的难点重新表述为：谁来给出足够好的 reward / judgment，以及这个 reward model 自身如何学习。

4 关键公式、算法与 readings 的连接

4.1 DPO: 本讲所有 recipe 的优化执行器

DPO reading 之所以重要，不是因为它单独就解决了 reasoning，而是因为它把“有偏好对就能训练”的这件事做得足够简单稳定。Weston 的系统几乎都默认用 DPO 或其变体做最后的优化，因此如果读者忽略 DPO，就会误以为这些方法的创新在于 loss 本身；实际上，创新更常在于偏好对是如何被构造出来的。

4.2 IRPO: 从 response preference 走向 reasoning preference

IRPO reading 正好补上了 lecture 的核心转折。它证明 reasoning task 上最大的痛点不是“没有 preference optimization”，而是“没有 reasoning-aware preference pairs”。一旦 final answer 可验证，

就可以把 CoT 也纳入 preference learning。

4.3 CoVe: 验证先于自信

CoVe reading 在本讲里承担的是“方法论校正”的角色。它提醒我们，面对 hallucination 问题时，最重要的未必是生成更多候选，而是把 verification workflow 明确写出来，让模型先问“我该如何检查自己”。

5 失败模式、边界条件与前后讲联系

5.1 本讲最重要的失败模式

1. 把 self-rewarding 理解成“模型随便给自己打分”。没有可解释 criteria 和种子 judge 数据时，这会迅速漂移。
2. 把 DPO 当成 reasoning 的全部。DPO 只是优化器；如果 winning/losing pairs 不携带真实 reasoning signal，它只会学到风格偏好。
3. 把 CoVe、S2A 当成更长 CoT。这些方法真正关心的是验证结构，而不是 token 数。
4. 在没有可验证 final answer 的任务上直接照搬 IRPO。一旦 reward grounding 变弱，reasoning preference 就会变得不可靠。
5. 忽略 judge 的学习问题。很多 agent pipeline 只训练 actor，却让 evaluator 原地踏步，这会导致系统上限卡死。

5.2 与前后讲的联系

第一讲讨论的是 inference-time compute 的分配：如何在推理时生成、筛选、修正 reasoning traces。本讲则把焦点移到 training time：如何让模型在训练中习得这些 reasoning / judging 行为。下一讲 Yu Su 会进一步把 reasoning 放进 language agent 框架，讨论 memory、planning、world model 等结构能力。这意味着课程主线正在从“如何让单个模型更会想”扩展到“如何让 agent 系统在外部环境中更会记、更会规划、更会行动”。

6 本章小结

Summary

- **Self-Rewarding** models can train themselves to get better – path to superhuman AI?
- Verifiable rewards help to train CoT for better reasoning (**Iterative Reasoning Preference Optimization, DeepSeek, O1**) & evaluation ability (**Thinking-LLM-as-judge**).
- Better judges (with CoT) can help train to *think* on non-verifiable tasks: **Thinking LLMs**.
- Models can even improve at **Meta-rewarding/reasoning** (*judging their judgements*).

Goal: An AI that "trains" itself as much as possible

- Creates new tasks to train on (challenges itself)
- Evaluates whether it gets them right ("self-rewarding")
- Updates itself based on what it understood

Research question: *can this help it become superhuman?*



图 13: 本讲 summary: reasoning learning 的升级路线, 最终都回到了 reward / judgment 是否足够可靠。

Weston 这讲最大的贡献, 是把“learning to reason”重新写成了“learning to generate better judgments”。其逻辑链条非常清楚:

- LLM 的 System 1 失败暴露了显式验证与去偏的必要性。
- 标准 RLHF 的 humans-in-the-loop 无法无限扩展。
- Self-Rewarding 让 actor 和 judge 一起成长。
- IRPO 说明 reasoning task 需要可验证 final answer 和 reasoning-specific preference pairs。
- Meta-Rewarding 与 EvalPlanner 进一步指出: judge 本身也是需要 reasoning 和训练的对象。

如果把本讲压缩成一句话, 那就是:

要让模型学会推理, 先得让系统学会更可靠地判断什么才算好的推理。

7 复习题

1. Weston 为什么把 hallucination、sycophancy、jailbreaking 都归为 System 1 failure?
2. CoVe 与单纯多采样 CoT 的核心差别是什么?
3. DPO 为什么只是 reasoning learning 的“执行器”, 而不是完整答案?
4. Self-Rewarding 模型为什么必须同时提升 acting 和 judging 两种能力?
5. IRPO 为什么需要 verifiable final answer 与 negative examples?

8 深入思考题

1. 如果一个任务没有明确 final answer，只能得到模糊的人类评分，你会怎样改写 IRPO 的 reward grounding?
2. Meta-Rewarding 是否会陷入 “judge judge judge” 的无限回归？系统需要在哪一层停下来并重新接入外部 ground truth?
3. 对真实 agent 环境而言，judge 最终应该主要依赖静态 benchmark、环境反馈，还是人类交互？为什么？

9 延伸阅读

- **DPO**: 理解为什么 KL-constrained RLHF 可以直接化为 preference optimization。
- **IRPO**: 理解 reasoning preference pair 的构造方式，以及为什么 NLL 项不能轻易拿掉。
- **CoVe**: 理解 verification planning 在 factuality 问题中的作用。
- 如继续深入本门课，可将本讲与第一讲的 verifier、Self-Debugging，以及后续 memory / planning 章节一起阅读，形成 “生成-评价-修正” 的统一视角。